


Logistic Regression- Part I

Lecture 6+7:



Outline



Example 1: Malignant or Benign

- Consider the example of recognizing whether a tumor in an input image is malignant or benign

Input:



Output:

Benign

0

Malignant

1

Can be
represented
as **integers**!


Can be represented as a **matrix** of pixels

Example 2: Spam or Not Spam

- As another example, consider the problem of detecting whether an email is a spam or not a spam

Input:

Email



Output:

Not Spam 0


Spam 1

Can be represented as **integers**!

Can be represented as a **vector** $\mathbf{x} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_d]$, with each component $\mathbf{x}_i$ corresponding to the presence ($\mathbf{x}_i = 1$) or absence ($\mathbf{x}_i = 0$) of a particular word (or *feature*) in the email

Regression vs. Classification


- What are the possible output values of the *linear regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x}$?



$\boldsymbol{\theta}$ is the *parameter vector*, that is, $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_m]$ (assuming $m + 1$ parameters) and $\mathbf{x}$ is the *feature vector*, that is, $\mathbf{x} = [\mathbf{x}_0, \mathbf{x}_1, \dots, \mathbf{x}_m]$ (assuming $m + 1$ features and $\mathbf{x}_0 = 1$ to account for the intercept term, namely, $\boldsymbol{\theta}_0$)


Regression vs. Classification

- What are the possible output values of the *linear regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x}$?
 - Real-valued outputs




Regression vs. Classification

- What are the possible output values of the *linear regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x}$?
 - Real-valued outputs



Regression vs. Classification

- What are the possible output values of the *linear regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x}$?
 - Real-valued outputs




We need a *discrete-valued* output (e.g., 1 or 0)

Regression vs. Classification

- How can we make the possible outputs of $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x}$ discrete-valued (as opposed to real-valued)?
 - By using an *activation function* (e.g., **sigmoid or logistic function**)

$$g(z) = \frac{1}{1 + e^{-z}}$$

$z \in \mathbb{R}$, but
 $g(z) \in [0,1]$



Assume a labeled example $(\mathbf{x}, y)$

If $y = 1$, we want $g(z) \approx 1$ (i.e. we want a correct prediction)


For this to happen, $z \gg 0$

Regression vs. Classification

- How can we make the possible outputs of $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x}$ discrete-valued (as opposed to real-valued)?
 - By using an *activation function* (e.g., **sigmoid or logistic function**)

$$g(z) = \frac{1}{1 + e^{-z}}$$

$z \in \mathbb{R}$, but
 $g(z) \in [0,1]$




Assume a labeled example $(\mathbf{x}, y)$

If $y = 0$, we want $g(z) \approx 0$ (i.e. we want a correct prediction)

For this to happen, $z \ll 0$


Regression vs. Classification

- How can we make the possible outputs of $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x}$ discrete-valued (as opposed to real-valued)?
 - By using an *activation function* (e.g., **sigmoid or logistic function**)




Regression vs. Classification

- How can we make the possible outputs of $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x}$ discrete-valued (as opposed to real-valued)?
 - By using an *activation function* (e.g., **sigmoid or logistic function**)



Regression vs. Classification


- How can we make the possible outputs of $\mathbf{h}_{\theta}(\mathbf{x}) = \theta^T \mathbf{x}$ discrete-valued (as opposed to real-valued)?
 - By using an *activation function* (e.g., **sigmoid or logistic function**)



So, What is Logistic Regression?

- Logistic regression is a machine learning algorithm that can be used to *classify* input data into discrete output (e.g., *input emails into spam or non-spam and tumour images into benign or malignant*)
 - **Note:** The word “regression” in the name does not mean that the algorithm is a regression algorithm (rather it is a classification algorithm)
- Major questions about logistic regression:
 - What is the *hypothesis function* (or *model*)?
 - What is the *cost function*?
 - How can we learn the parameters of the model?

Outline



The Logistic Regression Model

- What will be the output of the model $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x}$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Real-valued

- How can we make the output of $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})$ discrete?
 - By using the logistic function as follows:


$$\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x}) = \frac{1}{1 + e^{-\boldsymbol{\theta}^T \mathbf{x}}}$$

This is the logistic regression model or hypothesis function

- And then applying thresholding *after* learning the model to predict the output as follows

$$\begin{cases} \text{if } \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) < 0.5 \text{ predict } 0 \\ \text{if } \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) \geq 0.5 \text{ predict } 1 \end{cases}$$

Outline



Towards Identifying the Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Perhaps, by minimizing *Mean Squared Error (MSE)*. That is:

$$\text{minimize } \frac{1}{2n} \sum_{i=1}^n (y'^{(i)} - y^{(i)})^2$$

≡

$$\text{minimize}_{\boldsymbol{\theta}} \frac{1}{2n} \sum_{i=1}^n ((g(\boldsymbol{\theta}^T \mathbf{x}))^{(i)} - y^{(i)})^2$$

Cost function

$J(\boldsymbol{\theta})$

≡


$$\text{minimize}_{\boldsymbol{\theta}} \mathbf{J}(\boldsymbol{\theta})$$

Towards Identifying the Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Perhaps, by minimizing *Mean Squared Error (MSE)*. That is:

$$\text{minimize } \frac{1}{2n} \sum_{i=1}^n (y'^{(i)} - y^{(i)})^2$$

Easy for
Gradient
Descent
to locate!



$$\begin{aligned} &\equiv \\ &\text{minimize}_{\boldsymbol{\theta}} \frac{1}{2n} \sum_{i=1}^n ((g(\boldsymbol{\theta}^T \mathbf{x}))^{(i)} - y^{(i)})^2 \\ &\equiv \\ &\text{minimize}_{\boldsymbol{\theta}} \mathbf{J}(\boldsymbol{\theta}) \end{aligned}$$


Unfortunately, if we plot this cost function, it will turn out to be "**non-convex**"

Towards Identifying the Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Perhaps, by minimizing *Mean Squared Error (MSE)*. That is:

Gradient Descent might get stuck at a *local* min and fail to locate the *global* min!

Non-Convex



$$\text{minimize } \frac{1}{2n} \sum_{i=1}^n (y^{(i)} - y'^{(i)})^2$$

$\equiv$

$$\text{minimize}_{\boldsymbol{\theta}} \frac{1}{2n} \sum_{i=1}^n ((g(\boldsymbol{\theta}^T \mathbf{x}))^{(i)} - y^{(i)})^2$$

$\equiv$

$$\text{minimize}_{\boldsymbol{\theta}} \mathbf{J}(\boldsymbol{\theta})$$

Unfortunately, if we plot this cost function, it will turn out to be "**non-convex**"

Towards Identifying the Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Perhaps, by minimizing *Mean Squared Error (MSE)*. That is:

We need a cost function that is *convex*, hence, being more suitable for Gradient Descent

$$\text{minimize}_{\theta} \frac{1}{2n} \sum_{i=1}^n (y'^{(i)} - y^{(i)})^2$$

≡

$$\text{minimize}_{\theta} \frac{1}{2n} \sum_{i=1}^n ((g(\theta^T \mathbf{x}))^{(i)} - y^{(i)})^2$$

≡

$$\text{minimize}_{\theta} \mathbf{J}(\boldsymbol{\theta})$$

Unfortunately, if we plot this cost function, it will turn out to be "**non-convex**"

The Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Let us try a different cost function. That is:


$$\text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) = \begin{cases} -\log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$



The Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Let us try a different cost function. That is:


$$\text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) = \begin{cases} -\log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$



The Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Let us try a different cost function. That is:


$$\text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) = \begin{cases} -\log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$



The Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Let us try a different cost function. That is:

$$\text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) = \begin{cases} -\log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$




$$\text{If } y = 1 \begin{cases} \text{As } \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) \rightarrow 0, & -\log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) \rightarrow \infty \\ \text{As } \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) \rightarrow 1, & -\log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) \rightarrow 0 \text{ (i.e., cost } \rightarrow 0) \end{cases}$$

The Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Let us try a different cost function. That is:

$$\text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) = \begin{cases} -\log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$




This captures the intuition that if $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = 0$ (i.e., what we will predict is 0), but $y = 1$ (hence, we are mispredicting!), we shall penalize the learning algorithm by a very large cost!

The Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Let us try a different cost function. That is:


$$\text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) = \begin{cases} -\log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$



The Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Let us try a different cost function. That is:

$$\text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) = \begin{cases} -\log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$




The Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Let us try a different cost function. That is:

$$\text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) = \begin{cases} -\log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$

$$\text{If } y = 0 \begin{cases} \text{As } \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) \rightarrow 0, & -\log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) \rightarrow \infty \\ \text{As } \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) \rightarrow 1, & -\log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) \rightarrow \infty \end{cases} \quad (\text{i.e., cost} \rightarrow \infty)$$




The Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Let us try a different cost function. That is:

$$\text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) = \begin{cases} -\log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$

This captures the intuition that if $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = 1$ (i.e., what we will predict is 1), but $y = 0$ (i.e., we are mispredicting!), we shall penalize the learning algorithm by a very large cost!



The Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Let us try a different cost function. That is:

$$\text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) = \begin{cases} -\log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$

Equivalent To

$$\text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) = y \log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) - (1 - y) \log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}))$$

The Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Let us try a different cost function. That is:

$$\text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) = y \log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) - (1 - y) \log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}))$$


This function still assumes real-valued outputs for $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})$ (i.e., still entails *a regression problem*), while logistic regression should predict discrete values (i.e., logistic regression is a *classification problem*)

The Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - Let us try a different cost function. That is:

$$\text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) = y \log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) - (1 - y) \log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}))$$


We still need to apply to it the logistic function:

$$\mathbf{g}(z) = \frac{1}{1 + e^{-z}}$$

The Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - By minimizing the following cost function:


$$\begin{aligned} \text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) &= -y \log(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) - (1 - y) \log(1 - \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x})) \\ &= -y \log(\mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})) - (1 - y) \log(1 - \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})) \\ &= -y \log\left(\frac{1}{1 + e^{-\boldsymbol{\theta}^T \mathbf{x}}}\right) - (1 - y) \log\left(1 - \frac{1}{1 + e^{-\boldsymbol{\theta}^T \mathbf{x}}}\right) \end{aligned}$$

The Logistic Regression Cost Function

- How to learn a *logistic regression model* $\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{g}(\boldsymbol{\theta}^T \mathbf{x})$, where $\boldsymbol{\theta} = [\boldsymbol{\theta}_0, \dots, \boldsymbol{\theta}_m]$ and $\mathbf{x} = [\mathbf{x}_0, \dots, \mathbf{x}_m]$?
 - By minimizing the following cost function:

$$\text{Cost}(\mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}), y) = -y \log \left(\frac{1}{1 + e^{-\boldsymbol{\theta}^T \mathbf{x}}} \right) - (1 - y) \log \left(1 - \frac{1}{1 + e^{-\boldsymbol{\theta}^T \mathbf{x}}} \right)$$

If $y = 1$, we want $\boldsymbol{\theta}^T \mathbf{x} \gg 0$



If $y = 0$, we want $\boldsymbol{\theta}^T \mathbf{x} \ll 0$

